


European Football

(Known as soccer)
Justin O'Donnell, Brian Pascente, Carson Holscher,
Jacob Richardson

Basics of Soccer



Basics of Soccer Data

xG - Shot location (distance and angle to goal), body part used (foot, head, etc.), type of assist (cross, through ball, etc.), defensive pressure, and game context (e.g. counterattack, set-piece)

xA - the player who passed the ball before is credited with the xG of the shot that follows

xGBuildUp - If the player touches the ball before the shot is taken, then the xG is added to that player. This resets every time to other team gains possession

xGChain - Combination of xG, xA, and xGBuildup

Production per value (PPV)

- Our goal is to predict the value of soccer players based on all of their data from the season prior. This is to simulate when a team wants to buy another team's player and predict how good they will be.

$$\text{PPV (Midfielders)} = \frac{\text{Goals} \cdot \mathbf{0.4} + \text{Assists} \cdot \mathbf{0.4} + \text{Points Per Game} \cdot \mathbf{0.2}}{\text{Player Value}}$$

$$\text{PPV (Forwards)} = \frac{\text{Goals} \cdot \mathbf{0.4} + \text{Assists} \cdot \mathbf{0.3} + \text{Points Per Game} \cdot \mathbf{0.3}}{\text{Player Value}}$$

Literature Review

- Gained inspiration for what features and models to use, as many used Random Forests, Gradient Boosting, SVM, and even Neural Networks
- Many had features to be expected like age, height, market value
- Some predictors had weather conditions and injury/psychological evaluation
- One paper discussed player price optimization as a target, which was our main inspiration for our finalized target

Dataset

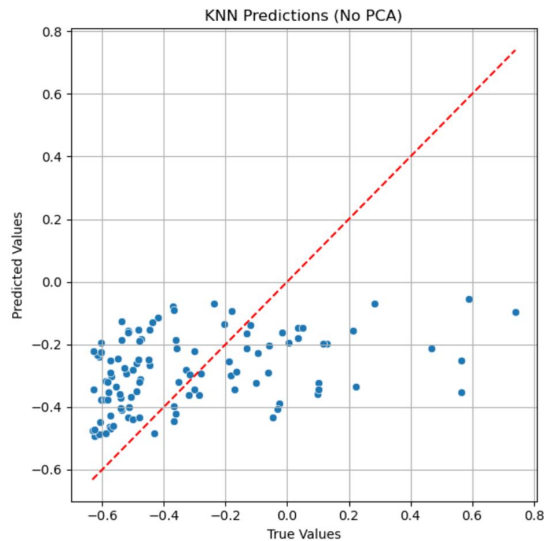
	minutes_played	goals	npg	assists	xG	xA	npxG	position_x	shots	key_passes	yellow_cards	red_cards
128	2423	0	0	2	0.886058	1.697511	0.886058	M S	21	22	8	0
106	2189	0	0	3	0.217996	1.706203	0.217996	D M S	8	17	2	0
331	996	2	2	0	1.234667	0.534926	1.234667	M S	7	8	2	0
167	2986	1	1	0	1.824859	1.692150	1.824859	D S	22	4	5	0
425	3017	0	0	2	0.134310	0.932408	0.134310	D S	8	12	4	0

xGBuildup	xGChain	market_value_in_eur	height_in_cm	age_in_months_2015	points_per_game	player_performance_valuation_standardized
8.497885	10.142531	800000.000000	177.000000	326.000000	1.314570	0.039118
5.533055	6.569234	100000.000000	180.000000	292.000000	1.213235	0.136234
1.861821	2.912522	1600000.000000	186.000000	344.000000	1.364035	-0.368572
5.403009	5.606335	400000.000000	193.000000	351.000000	1.967105	-0.075814
9.725240	10.318017	2000000.000000	182.000000	253.000000	1.157895	-0.414862

Model Performance (R^2)

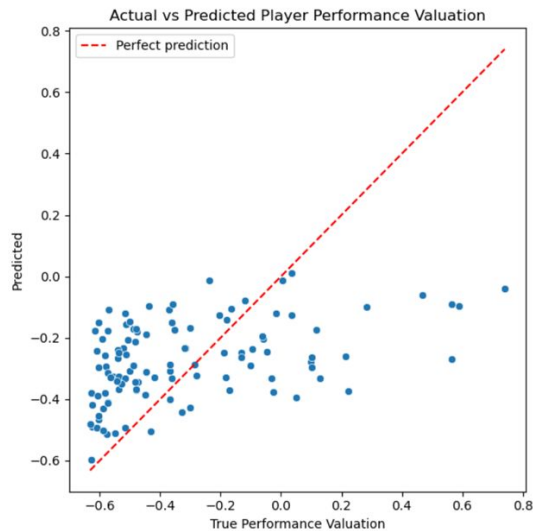
Gradient Boost	0.5018265603
KNN	0.2747
Ridge	0.10
Forest	0.11721429293514218
Lasso	0.06

KNN + PCA and Feature Engineering




KNN (K=23)

Test $R^2 = 0.114$



KNN (PCA, K=23)

Test $R^2 = 0.169$



KNN (PCA+FE, K=15)

Test $R^2 = 0.275$

Gradient Boost


```
param_grid = {  
    'n_estimators': [100, 200, 300, 500],  
    'learning_rate': [0.01, 0.05, 0.1, 0.2],  
    'max_depth': [3, 4, 5, 6, 7]  
}
```

Best Parameters: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 300}


MSE: 0.0296159665

MAE: 0.1215287545

Feature Importance:




Actual vs Predicted Values



Demo Example 1:

```
data = pd.DataFrame([
    {
        'minutes_played': 2823,
        'goals': 6,
        'npg': 6,
        'assists': 7,
        'xG': 2.794280,
        'xA': 5.305932,
        'npxG': 2.794280,
        'position_x': 'F M S',
        'shots': 65,
        'key_passes': 65,
        'yellow_cards': 7,
        'red_cards': 1,
        'xGBuildup': 4.623076,
        'xGChain': 8.855647,
        'market_value_in_eur': 300000.0,
        'height_in_cm': 173.0,
        'age_in_months_2015': 345.0,
        'points_per_game': 1.184211
    }
])
```



Predicted performance valuation: 0.2902878596


Jason Puncheon

Good Player

Expected: 0.258462

Demo Example 2:

```
data = [{  
    'minutes_played': 2682,  
    'goals': 12,  
    'npg': 12,  
    'assists': 7,  
    'xG': 9.096988,  
    'xA': 10.388413,  
    'npxG': 9.096988,  
    'position_x': 'M',  
    'shots': 66,  
    'key_passes': 92,  
    'yellow_cards': 8,  
    'red_cards': 0,  
    'xGBuildup': 16.633573,  
    'xGChain': 29.144278,  
    'market_value_in_eur': 50000.0,  
    'height_in_cm': 178.0,  
    'age_in_months_2015': 341.0,  
    'points_per_game': 1.927632  
}]
```




Predicted performance valuation: 0.4036456920

David Silva
Very Good Player
Expected: 11.3

Demo Example 3:

```
data = [{  
  'minutes_played': 1092,  
  'goals': 0,  
  'npg': 0,  
  'assists': 1,  
  'xG': 0.106066,  
  'xA': 0.344920,  
  'npxG': 0.106066,  
  'position_x': 'D S',  
  'shots': 7,  
  'key_passes': 8,  
  'yellow_cards': 0,  
  'red_cards': 0,  
  'xGBuildup': 2.098642,  
  'xGChain': 2.395614,  
  'market_value_in_eur': 1800000.0,  
  'height_in_cm': 179.0,  
  'age_in_months_2015': 267.0,  
  'points_per_game': 1.169173  
}]
```



Predicted performance valuation: -0.4100520727

Bad Player

Massadio Haidara

Expected: -0.461564

Demo 4: Random player

Future Work

- Probably look into aggregating model results. The author of one paper that we read used a combination of random forest regression, support vectors & gradient boosting.
- Remove the bench players when evaluating higher performing players.
- Fix the joins so that we don't unnecessarily delete a good chunk of our data, due to team names.
- Procure a more complete dataset (such as the one used for FIFA)

Lessons

- Overfitting is very easy to do in the real world
- There's no perfect metric for what makes a good soccer player. Your best bet for figuring it out is common sense, but that'll only get you so far.
- A few bad assumptions can severely affect your model's performance

Questions?